


Artificial intelligence and chemistry

Tony Bonnaire

Image generated with Midjourney « A representation of deep learning merging with chemistry »

 **Given by:** Tony Bonnaire (+ **Aurora González Sanz** for the project)

 **Format:** Lectures + hands-on sessions then data challenge (in chemistry)

 **Exam:** Paper analysis (30%) + MCQ (30%) + oral presentation for the challenge (40%)

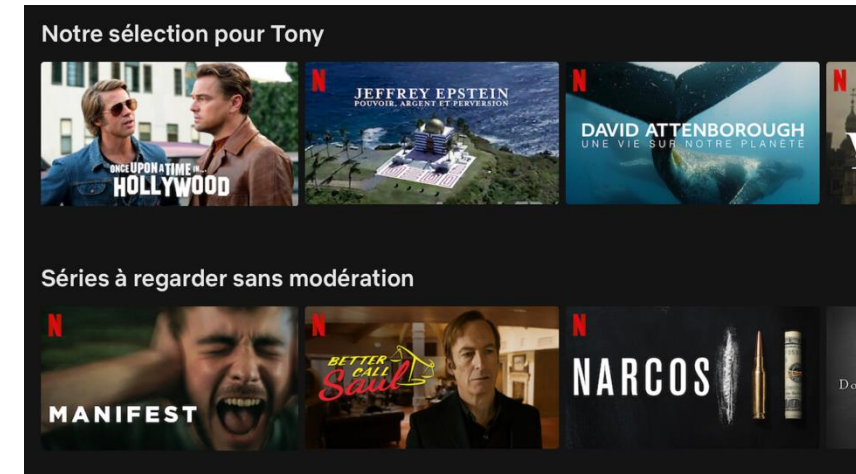
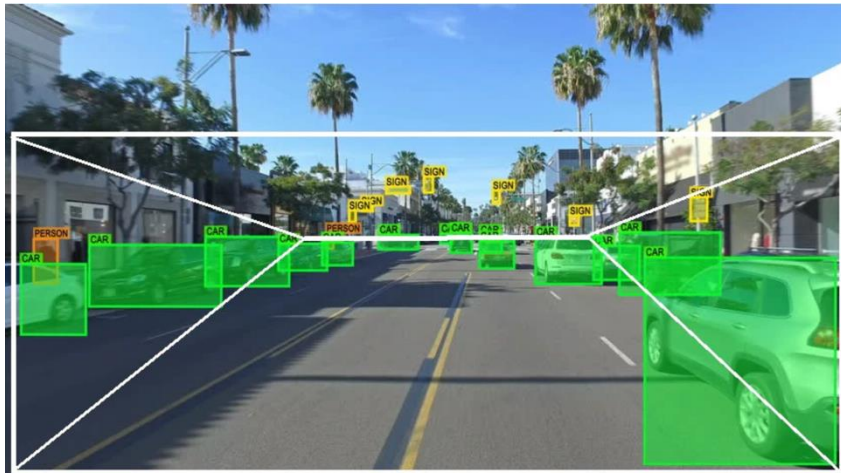
 **Aim:** Introduce you to the basics of ML principles and carry out a project

Some references:

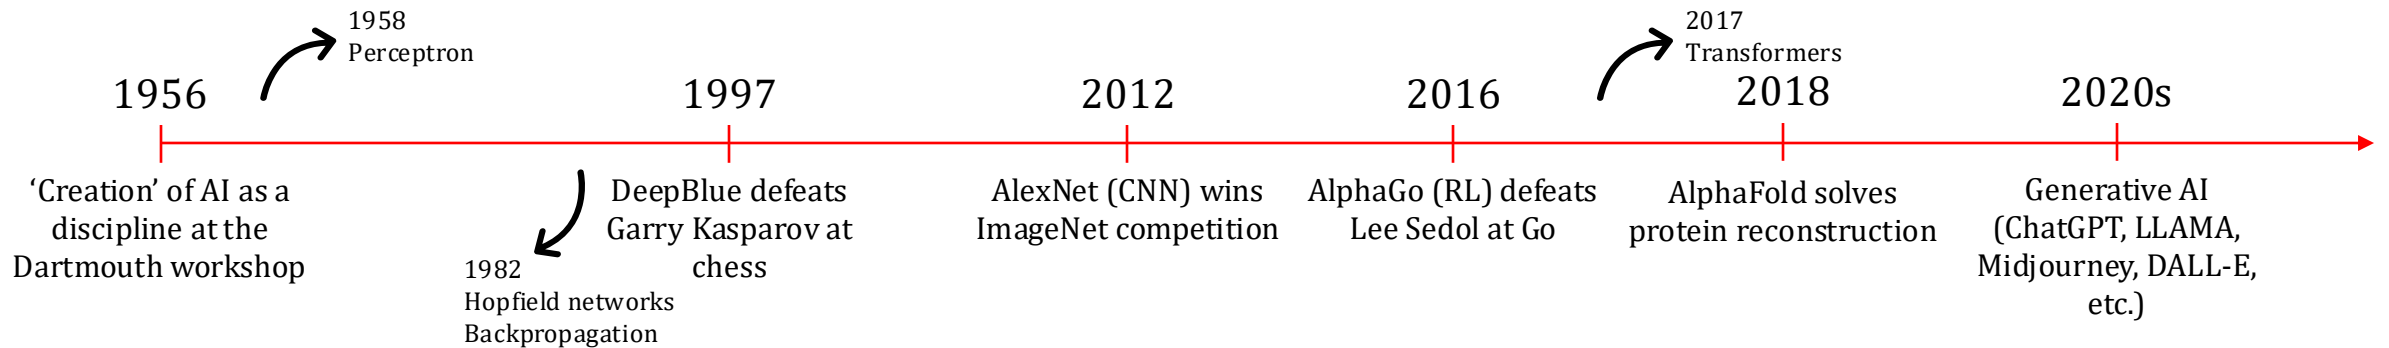
- [Deep Learning: Foundations and Concepts](#), Bishop & Bishop, 2023,
- [Deep Learning](#), Goodfellow et al., 2016,
- [Deep Learning with Python](#), Chollet, 2016,
- [Learning Theory from First Principles](#), Francis Bach, 2024,
- <https://challengedata.ens.fr>: a bank of data science challenges to apply all the things we will learn in this course

AI goal


Design **systems** capable of performing complex tasks requiring *intelligence* (i.e. using reasoning, perception or language) **to take decisions** and **make predictions**.



Some (selected) AI breakthroughs



AI in science



2025 not over yet!

80 papers a day in average in 2024 (!)

Some scientific applications

Healthcare

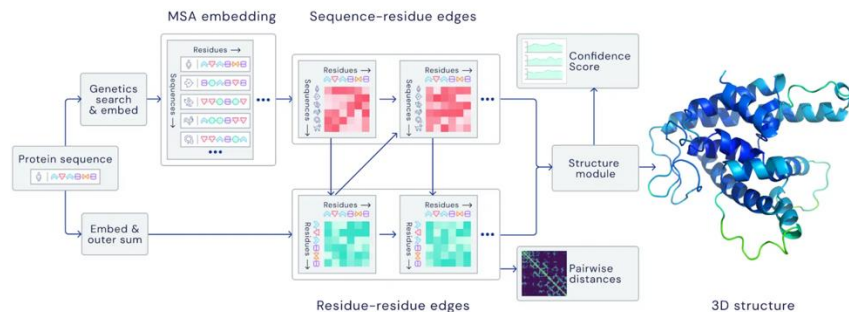
- Drug discovery
- Protein structure reconstruction

Astrophysics and cosmology


- Galaxy deblending
- Image restoration
- Source separation

Theoretical physics


- Study phase transitions
- Discover experiments and equations



[Jumper et al., 2021](#)



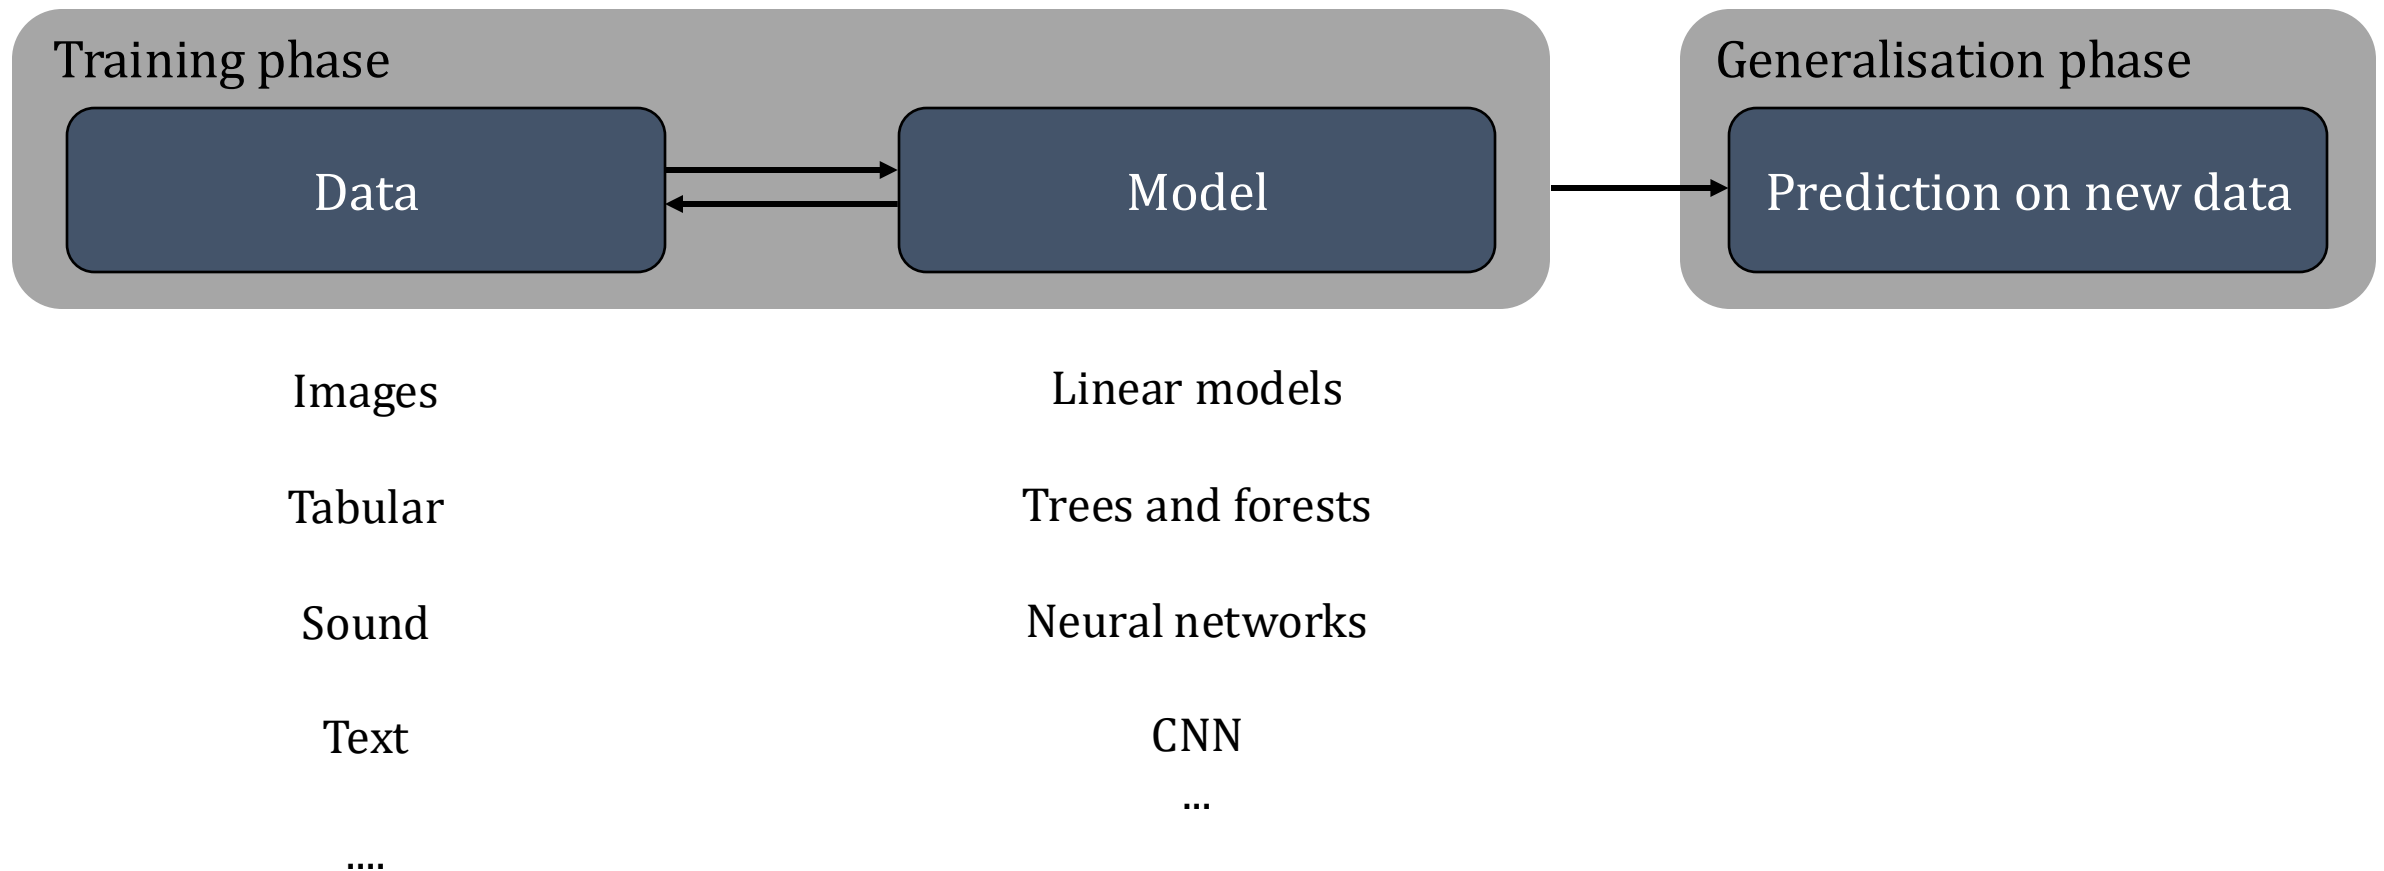
[Medeiros et al., 2023](#)



[Van Nieuwenburg et al., 2017](#)

... And many more (climate forecast, fraud detection in cybersecurity, binding energies in quantum chemistry)

Machine Learning came as a solution to design intelligent systems, replacing handcrafted decision rules by **learnt rules** using **training data** and **optimisation** of **parameterised models**.






What is “learning”? An example


6



Images of a “cat” or “dog”

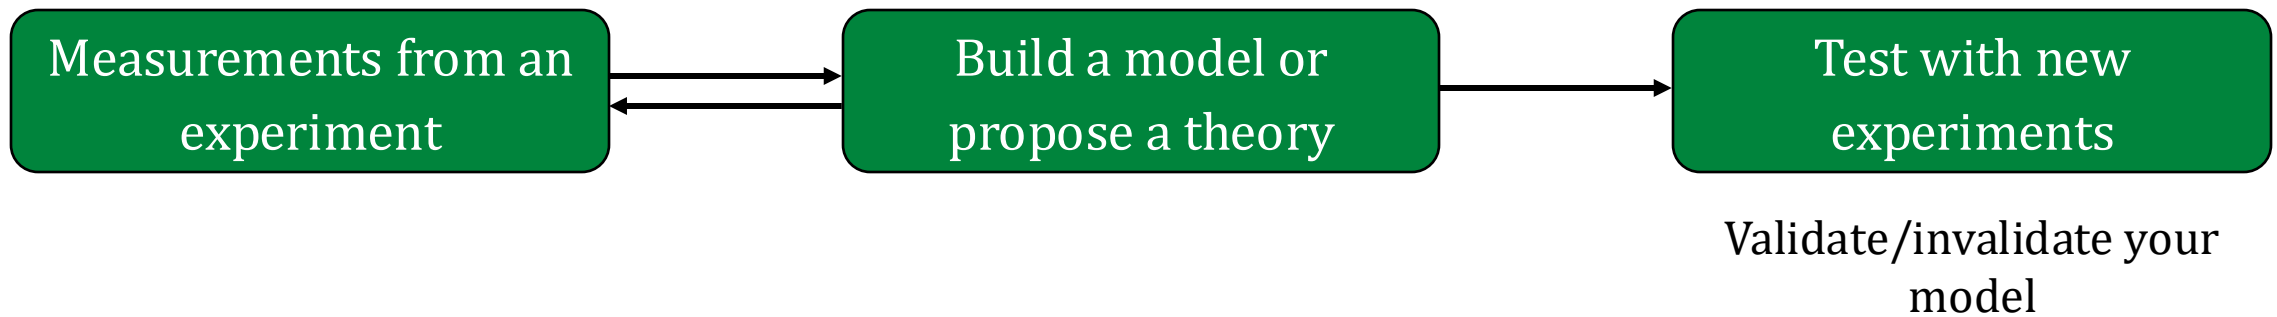


“cat” or “dog” ?

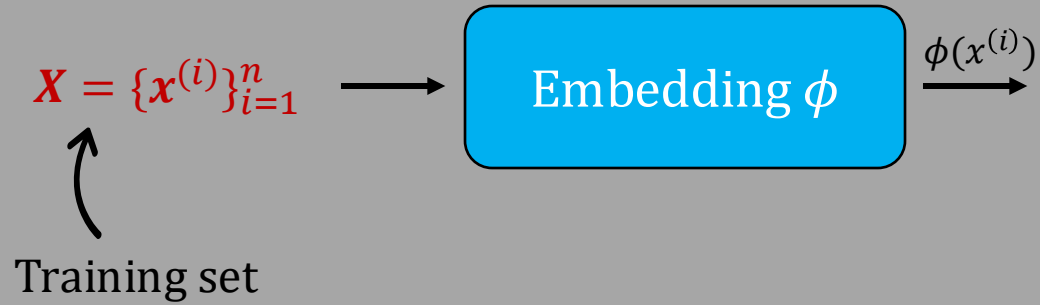


...In fact, all this is close to what you know!

The scientific method



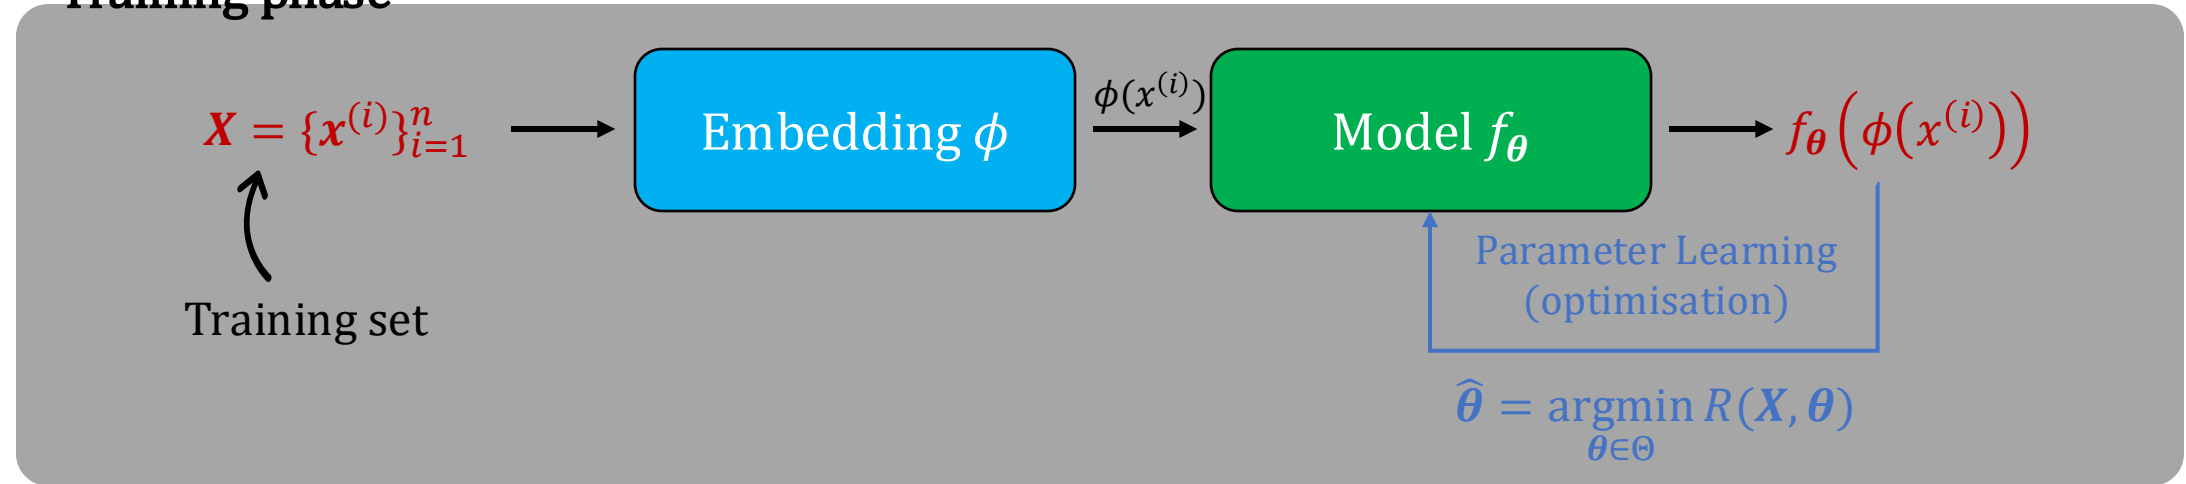
Training phase



1. Data X are **unstructured**, sometimes **noisy** and **unprocessed** like pixels of an image or sequence of characters or words.
2. The embedding $\phi(x^{(i)})$ is a **structured, numerical** vector representation of the data whose elements are **meaningful features**. It depends on the data and the purpose. It can be **handcrafted or learnt**.

Finding a good embedding is a central part of ML: it eases the problem by preserving the essential structure of the data that matters for the task.

Training phase

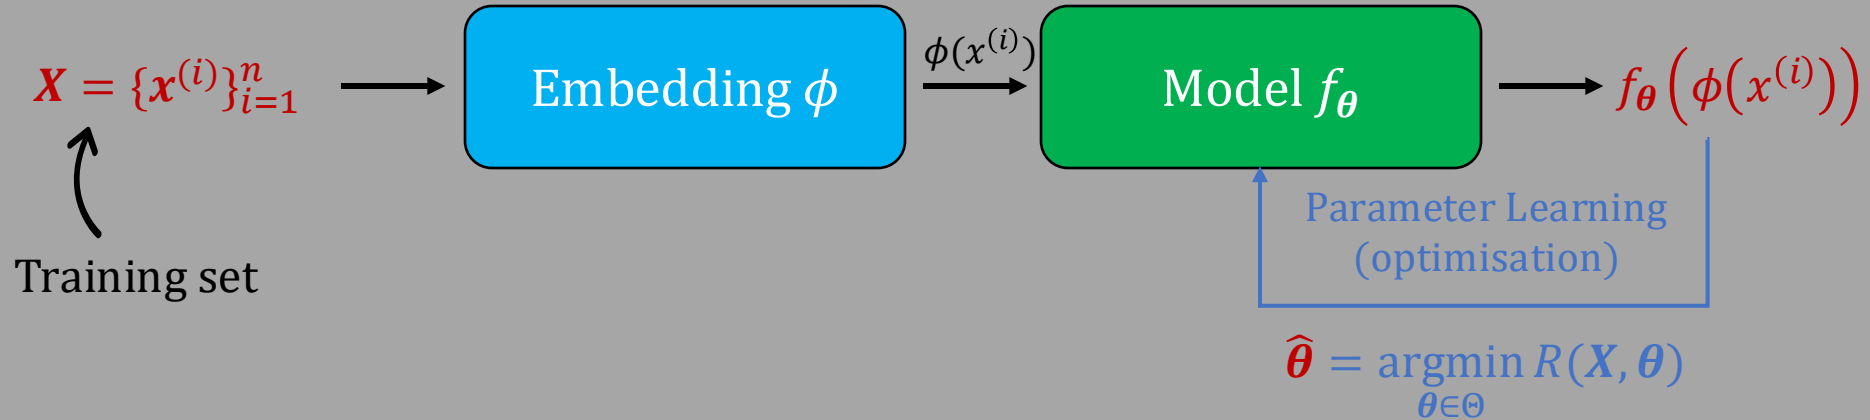


Some notations and terminologies:

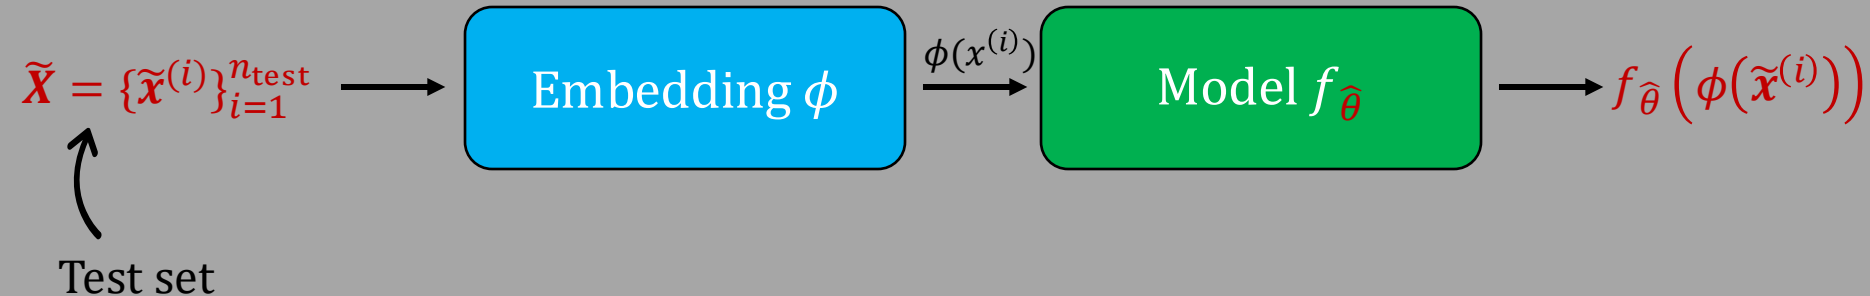
- $x^{(i)} \in \mathbb{R}^d$ is *one training data* (there are n of them),
- $\phi(x^{(i)}) \in \mathbb{R}^{d'}$ is an embedding of $x^{(i)}$ sometimes called *feature vector*,
- $\theta \in \Theta \subset \mathbb{R}^p$ are the *parameters* of the model,
- $R(X, \theta)$ is the *risk* and measures the error of the model with parameters θ on data X .

At the end of the training procedure, we have a model $f_{\hat{\theta}}$ committing an error of $R_{\text{train}} = R(X, \hat{\theta})$ on the training set.

Training phase



Generalisation phase



Using the test set, we can evaluate the test error $R_{\text{test}} = R(\tilde{X}, \hat{\theta})$ and compare it to R_{train} to detect **generalisation issues** (overfitting or underfitting).

Supervised learning

- Training data are actually X and Y coming as pairs

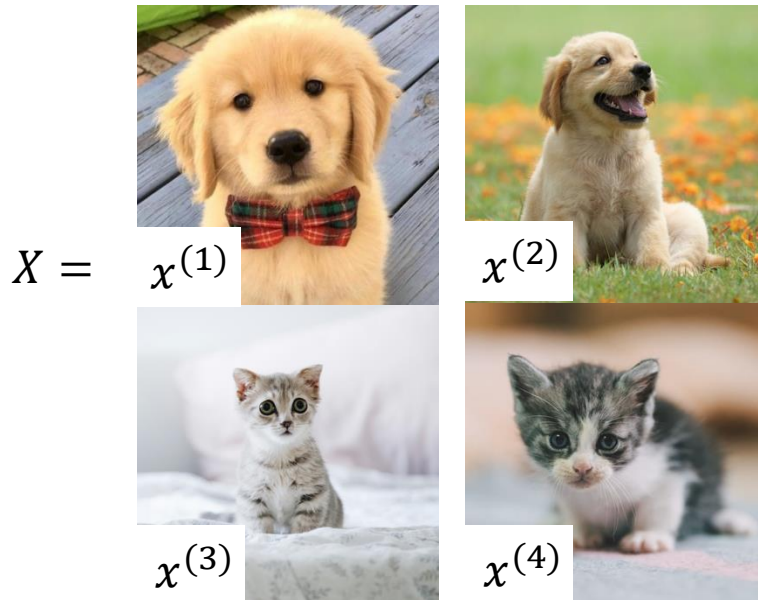
$$X = \{(\mathbf{x}^{(i)}, \mathbf{y}^{(i)})\}_{i=1}^n, \quad (\mathbf{x}^{(i)}, \mathbf{y}^{(i)}) \in \mathbb{X} \times \mathbb{Y}$$

- If \mathbb{Y} is continuous, then the task is called **regression**, and if \mathbb{Y} is discrete, then it is a **classification** problem.



$\mathbf{x}^{(i)}$ is the i th **data vector** of the training base, and $\mathbf{y}^{(i)}$ is called the **target (or predicted) variable**

Example: Determine if an image encodes a cat or a dog (called a **classification** task)



$$Y = \{1, 1, 0, 0\}$$

$$f_{\theta}(\mathbf{x}^{(i)}) = \hat{\mathbf{y}}^{(i)}$$



$$f_{\hat{\theta}}$$

$$\hat{\theta} = \underset{\theta}{\operatorname{argmin}} R(X, \theta)$$

Training data

Model

Optimisation

Supervised learning

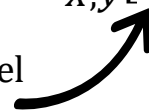
- Training data are actually X and Y coming as pairs

$$\mathbf{X} = \{(\mathbf{x}^{(i)}, \mathbf{y}^{(i)})\}_{i=1}^n, \quad (\mathbf{x}^{(i)}, \mathbf{y}^{(i)}) \in \mathbb{X} \times \mathbb{Y}$$

- If \mathbb{Y} is continuous, then the task is called **regression**, and if \mathbb{Y} is discrete, then it is a **classification** problem.
- Ideally, we would like to minimise the **expected risk**, i.e. the **expected value of a loss function** $\ell(y, \hat{y})$

$$R(\mathbf{X}, \boldsymbol{\theta}) = \mathbb{E}_{\mathbf{X}, \mathbf{y}}[\ell(\mathbf{y}, \hat{\mathbf{y}})]$$

Loss function: measures how bad your model
is on a single example



However, we do not know $p(\mathbf{X}, \mathbf{y})$ so in practice we rely on the **empirical risk** instead

$$\hat{R}(\mathbf{X}, \boldsymbol{\theta}) = \frac{1}{n} \sum_{i=1}^n \ell(\mathbf{y}^{(i)}, \hat{\mathbf{y}}^{(i)}).$$

Supervised learning

- Training data are actually X and Y coming as pairs

$$\mathbf{X} = \{(\mathbf{x}^{(i)}, \mathbf{y}^{(i)})\}_{i=1}^n, \quad (\mathbf{x}^{(i)}, \mathbf{y}^{(i)}) \in \mathbb{X} \times \mathbb{Y}$$

- If \mathbb{Y} is continuous, then the task is called regression, and if \mathbb{Y} is discrete, then it is a classification problem.

Examples of
tasks

Classification

Regression

Timeseries prediction

Segmentation

Examples of
models

Artificial Neural network

Random forest

Linear regression

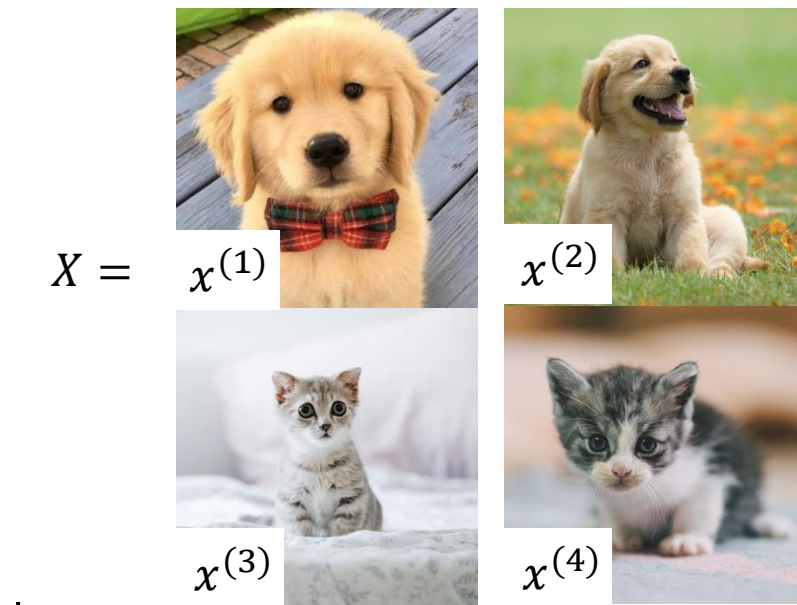
Logistic regression

Naïve Bayes

Nearest neighbours

Unsupervised learning

- Training data are the set of $\mathbf{x}^{(i)}$'s only; no known results to predict
- In unsupervised learning, one seeks **patterns or structures** in X without prior labels
- Usually boils down to model the probability distribution of the dataset



Training data

Example: Generate new images of cats and dogs (called a **generation** task)

$$f_{\theta}(\mathbf{x}) = p_{\theta}(\mathbf{x}) \longrightarrow f_{\hat{\theta}}$$
$$\hat{\theta} = \underset{\theta}{\operatorname{argmin}} R(\mathbf{X}, \theta)$$

such that $p_{\theta}(\mathbf{x}) \approx p(\mathbf{x})$

Model

Optimisation

Unsupervised learning

- Training data are the set of $\mathbf{x}^{(i)}$'s only; no known results to predict
- In unsupervised learning, one seeks **patterns or structures** in X without prior labels
- Usually boils down to model the probability distribution of the dataset

Examples of
tasks

Clustering

Data augmentation

Dimensionality reduction

Sampling

Examples of
models

Autoencoder

Boltzmann Machine

Diffusion models

Gaussian mixture model

Generative Adversarial network

Reinforcement learning

- The philosophy is different: the model does not try to “imitate” like in supervised learning nor to find patterns but “tries” things
- It is based on an **agent** interacting with an **environment**
- The agent tries to find the best possible sequence of states and actions to **maximise a reward**

Examples of tasks

Game theory

Robotics


Autonomous driving

Examples of models

Markov decision processes

Q-networks

Deep policy gradient



Not discussed in this course, but a very good reference is [Reinforcement Learning – An introduction, Sutton and Barto, 2018](#)



$d \approx 10^6$

- To sample a $[0,1]^d$ space with a shortest distance to a test point at most ϵ , we need $n_{\text{train}} \geq \epsilon^{-d} = e^{-d \log \epsilon}$
- $d \approx 80$ requires more samples than the number of atoms in the universe


1-nearest neighbour

- A simple classification rule is for instance associating to a data the label of its closest neighbour in the d -dimensional space.

$$\hat{y} = y^{(m)} \text{ with } m = \operatorname{argmin}_i \|x - x^{(i)}\|_2^2$$

- In this case $R_{\text{train}} = 0$ but R_{test} is very large! Why?

Curse of dimensionality





Traditional methods typically break down in high-dimensional spaces (**curse of dimensionality**) and it is impossible to design handcrafted decision rules for complex tasks.



IF THERE IS ONLY ONE THING TO REMEMBER FROM THIS CLASS

The **curse of dimensionality** is the **central problem of machine learning**. To fight it, ML relies on **prior information** about the problem:

- **Reduce the dimensionality**: select a subset of meaningful features (or their interactions) through appropriate embeddings.
- Exploit **structures** in the data (invariances, sparsity, long/short range correlations, etc.) to define the model,
- **Penalise** complex models leading to poor generalisation performances using regularisation.