


IV – Neural networks

Contents:

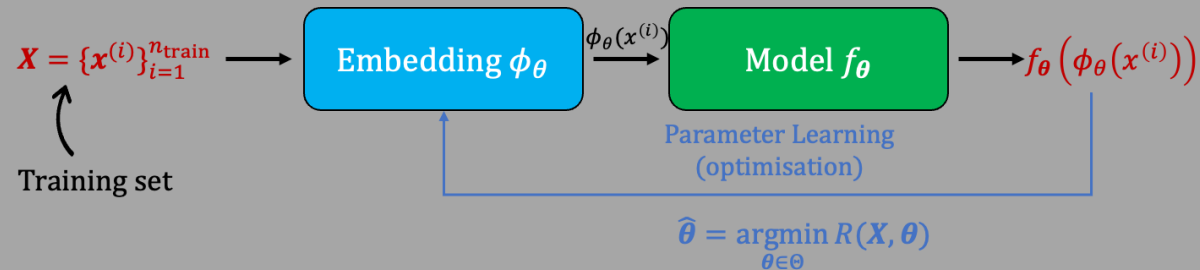
- *Neurons, activation functions*
- *Learning features with feed-forward neural networks*
- *Convolutional neural networks: invariances, convolutions, pooling*



Artificial neural networks

- How can we handcraft features $\phi(x)$ for complex problems like real-life image classification allowing the use of simple linear decision models?
- The idea of neural network is to parameterise the embedding $\phi_\theta(x)$ and find the basis linearising the problem.**


Training phase




Images from the ImageNet dataset

Artificial neural networks

- Building block of neural networks: the **neuron**
- Made of three operations: it first multiplies the input by a **weight**, then adds a **bias**, and finally applies an **activation function** s to the result
- If s is the identity, you recognize the linear regression model with $\theta_0 = b$ and $\theta_1 = w$. To represent non-linear functions, s must be non-linear




ReLU




$$s(x) = \max(0, x)$$

GeLU



$$s(x) = \frac{x}{2} \left(1 + \operatorname{erf} \left(\frac{x}{\sqrt{2}} \right) \right)$$

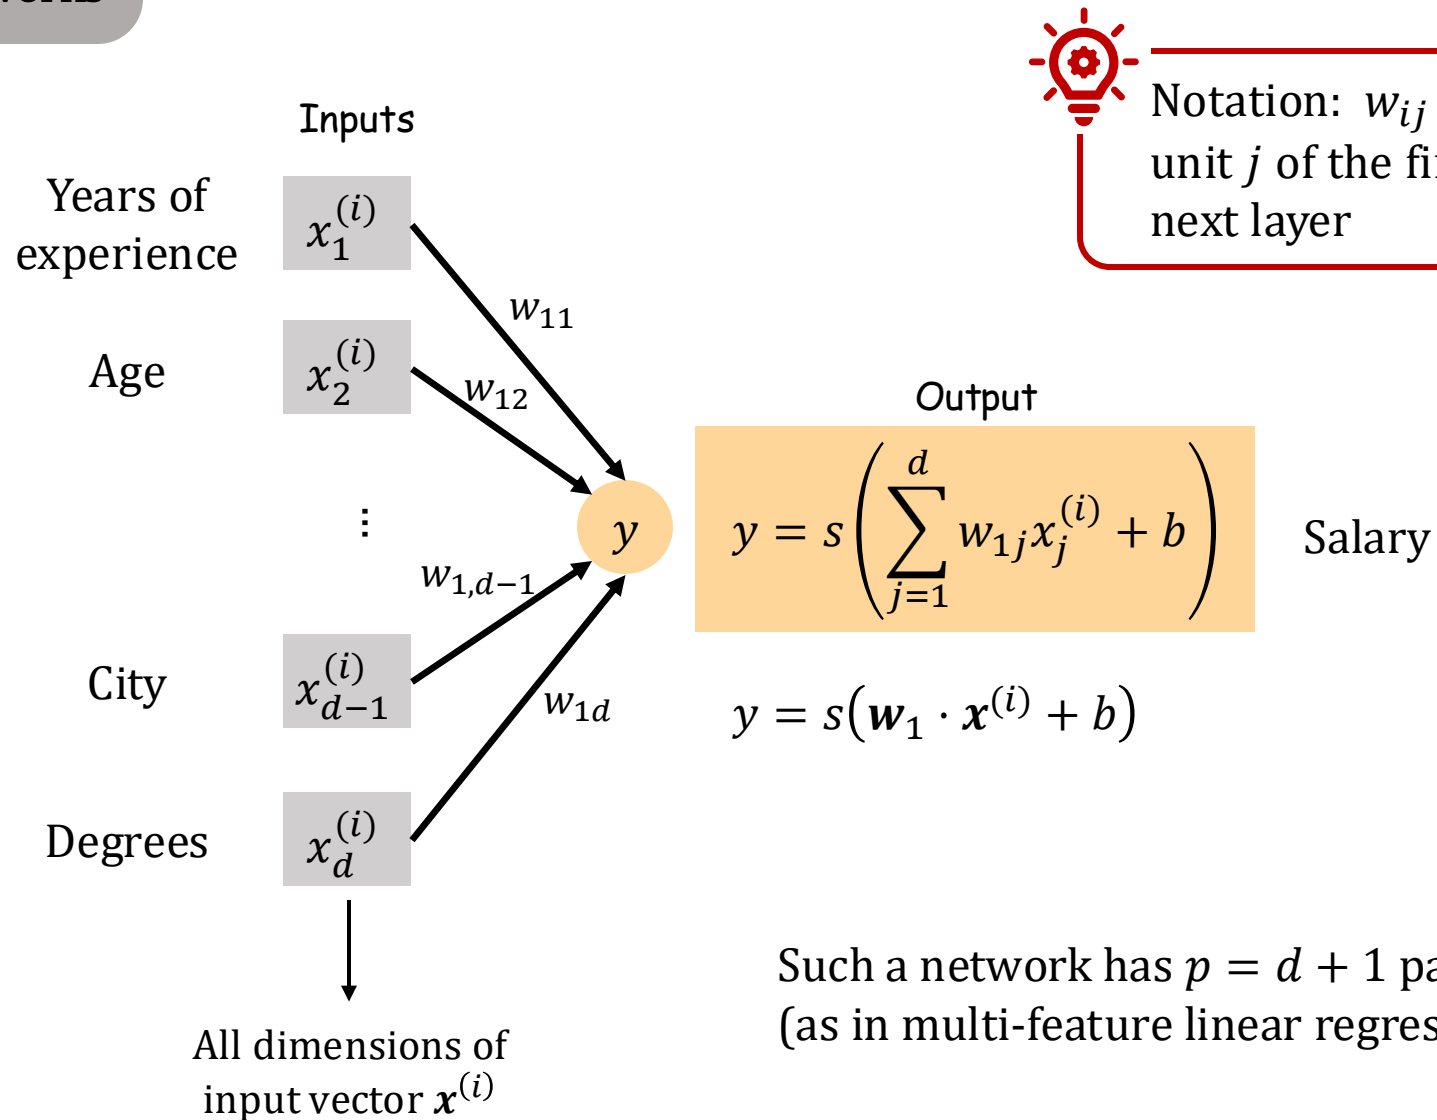
Sigmoid



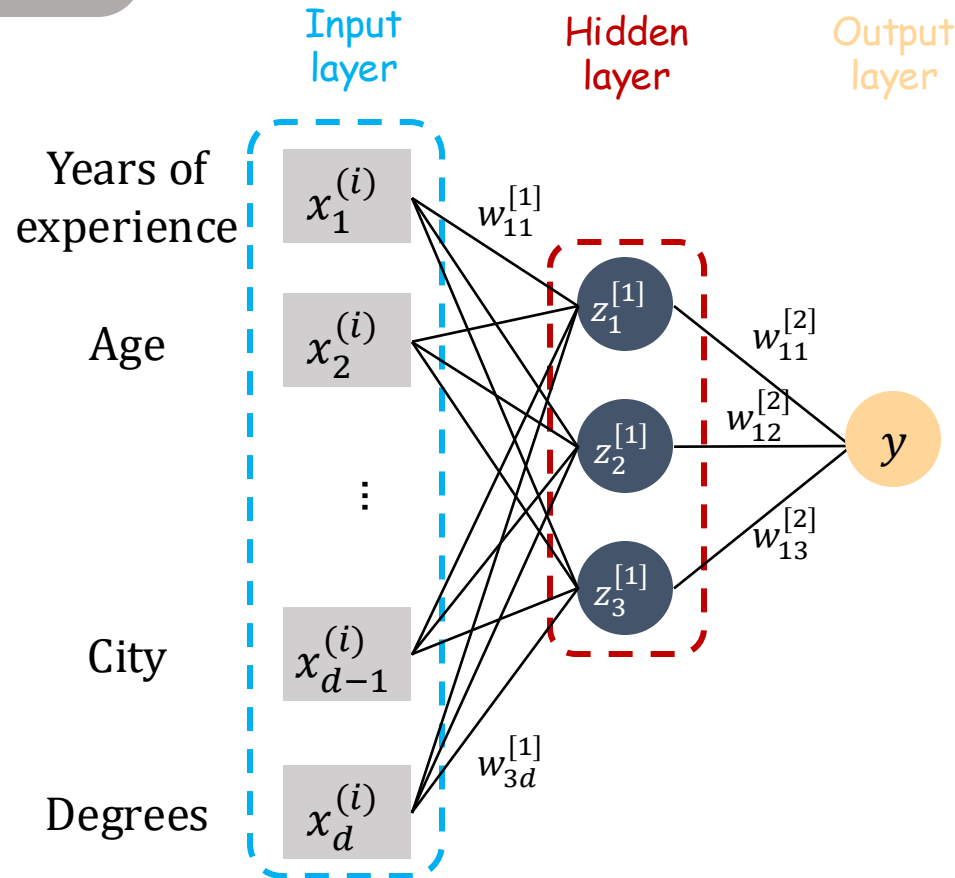
$$s(x) = \frac{1}{1 + \exp(-x)}$$

- This model is motivated by biological neurons and the activation function mimics the activation or inhibition through **non-linear (and differentiable) functions**
- Can take different forms but some examples include the **ReLU or sigmoid functions**

Artificial neural networks



Artificial neural networks



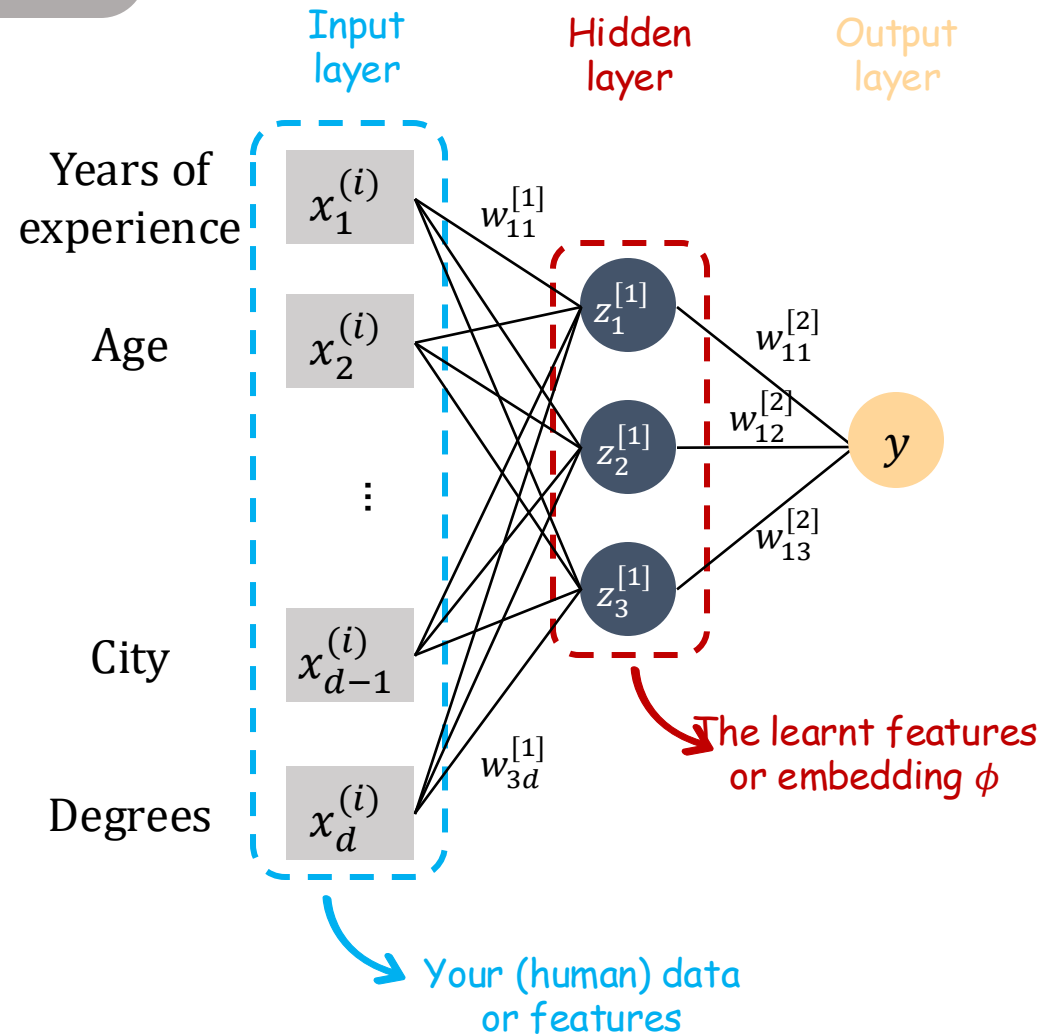
- Units in a same layer do not interact
- Data go from input to output in a **feed-forward way**
- The width of the output layer corresponds to the number of classes/values that you want to predict
- The **activation** of the unit j in layer one is given by

$$z_j^{[1]} = s \left(\underset{\substack{\text{Weights of} \\ \text{layer 1 for unit} \\ j \text{ (vectorized)}}}{w_j^{[1]}} \cdot \underset{\substack{\text{Unit } j}}{x^{(i)}} + \underset{\substack{\text{Bias of the} \\ \text{unit } j \text{ in} \\ \text{layer 1}}}{b_j^{[1]}}} \right)$$

- We also define the **pre-activation**

$$u_j^{[1]} = w_j^{[1]} \cdot x + b_j^{[1]}$$

Artificial neural networks



Let's have a look at the output of this network

$$y = s \left(\sum_{j=1}^d w_{1j}^{[2]} z_j^{[1]} + b^{[2]} \right)$$

$\{z_j\}_{j=1,\dots,3}$ act as **new features** to predict y

Artificial neural networks


- A fully-connected neural network with d inputs, L hidden layers of width h_1, h_2, \dots, h_L and an output layer of size q
- The j^{th} output is computed as

$$\mathbf{y} = s^{[L+1]}(\mathbf{W}^{[L+1]}\mathbf{z}^{[L]}) \quad \text{with } \mathbf{z}^{[L]} = [1, z_1^{[L]}, z_2^{[L]}, \dots, z_{h_L}^{[L]}]^T$$

$$\mathbf{y} = s^{[L+1]} \left(\mathbf{W}^{[L+1]} s^{[L]} \left(\mathbf{W}^{[L]} \dots s^{[1]}(\mathbf{W}^{[1]} \mathbf{x}) \right) \right)$$


$$\mathbf{w}_j^{[l]} = [b_j^{[l]}, w_{1j}^{[l]}, w_{2j}^{[l]}, \dots, w_{h_{l-1}j}^{[l]}]$$

$$\mathbf{W}^{[l]} = [\mathbf{w}_1^{[l]}, \mathbf{w}_2^{[l]}, \dots, \mathbf{w}_{h_{l-1}}^{[l]}]^T \in \mathbb{R}^{h_l \times (h_{l-1} + 1)}$$



- In the end, a **neural network** is a, as other models, a function $f_{\theta}: X \rightarrow Y$ of some parameters ($\theta = \text{weights and biases}$)
- f_{θ} is a composition of non-linear function when s is non-linear, allowing to build non-linear estimators
- The parameters of the model are obtained by minimising the empirical risk (**cross-entropy for classification, MSE for regression**)

Artificial neural networks



How many parameters does this network have? (don't forget biases!)

- + Both regression and classification, learns features, good performances when enough data
- Need a lot of data to train, subject to overfitting, uninterpretable, targeted-purpose architectures usually work better

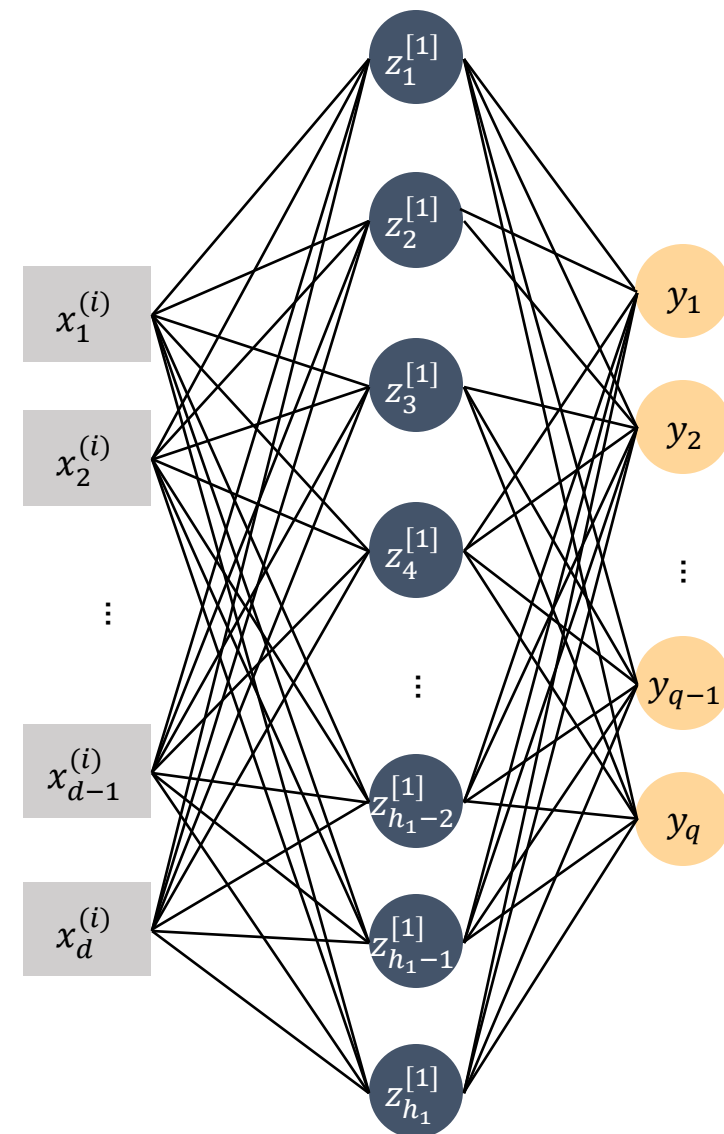
Artificial neural networks

- Single-layer neural networks are universal approximation (see [Cybenko 1989](#))

Theorem (informal)



*A fully-connected neural network with a single hidden layer ($L = 1$) with enough neurons (h_1 large) can fit **any arbitrary smooth function**.*

- In practice, this does **not** help: the width may need to scale exponentially with the function complexity.
- Empirically, **deep networks** are more efficient than wide ones.





Artificial neural networks



1 hidden layer of 1 unit



1 hidden layer of 100 units



2 hidden layers of 10 units



Linear models

Principles of learning

Trees and ensembling

Neural networks

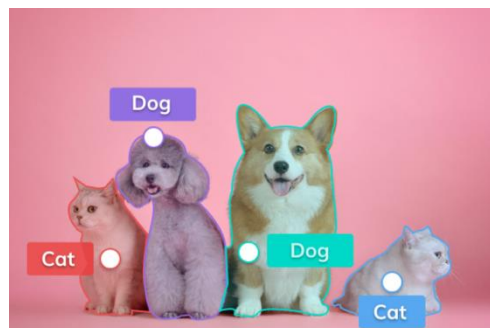
Risk optimization

Image classification




Dog?

Instance segmentation



Object detection



Caption generation



“A dog and a little boy playing with a basketball in the grass”

Image classification



Dog?

12M pixels

$$x^{(i)} = [x_1^{(i)}, x_2^{(i)}, \dots, x_{12192768}^{(i)}]$$

$$d = 12192768$$

Why going beyond fully-connected neural network?

- Consider the simplest task with **classification**
- Standard images taken by a current smartphone are of size $4032 \times 3024 = 12\text{M}$ pixels
- A 2 layer-FCNN with just 100 (!) units has $> 1\text{B}$ parameters




Image classification



A dog

Translation



Still a dog!

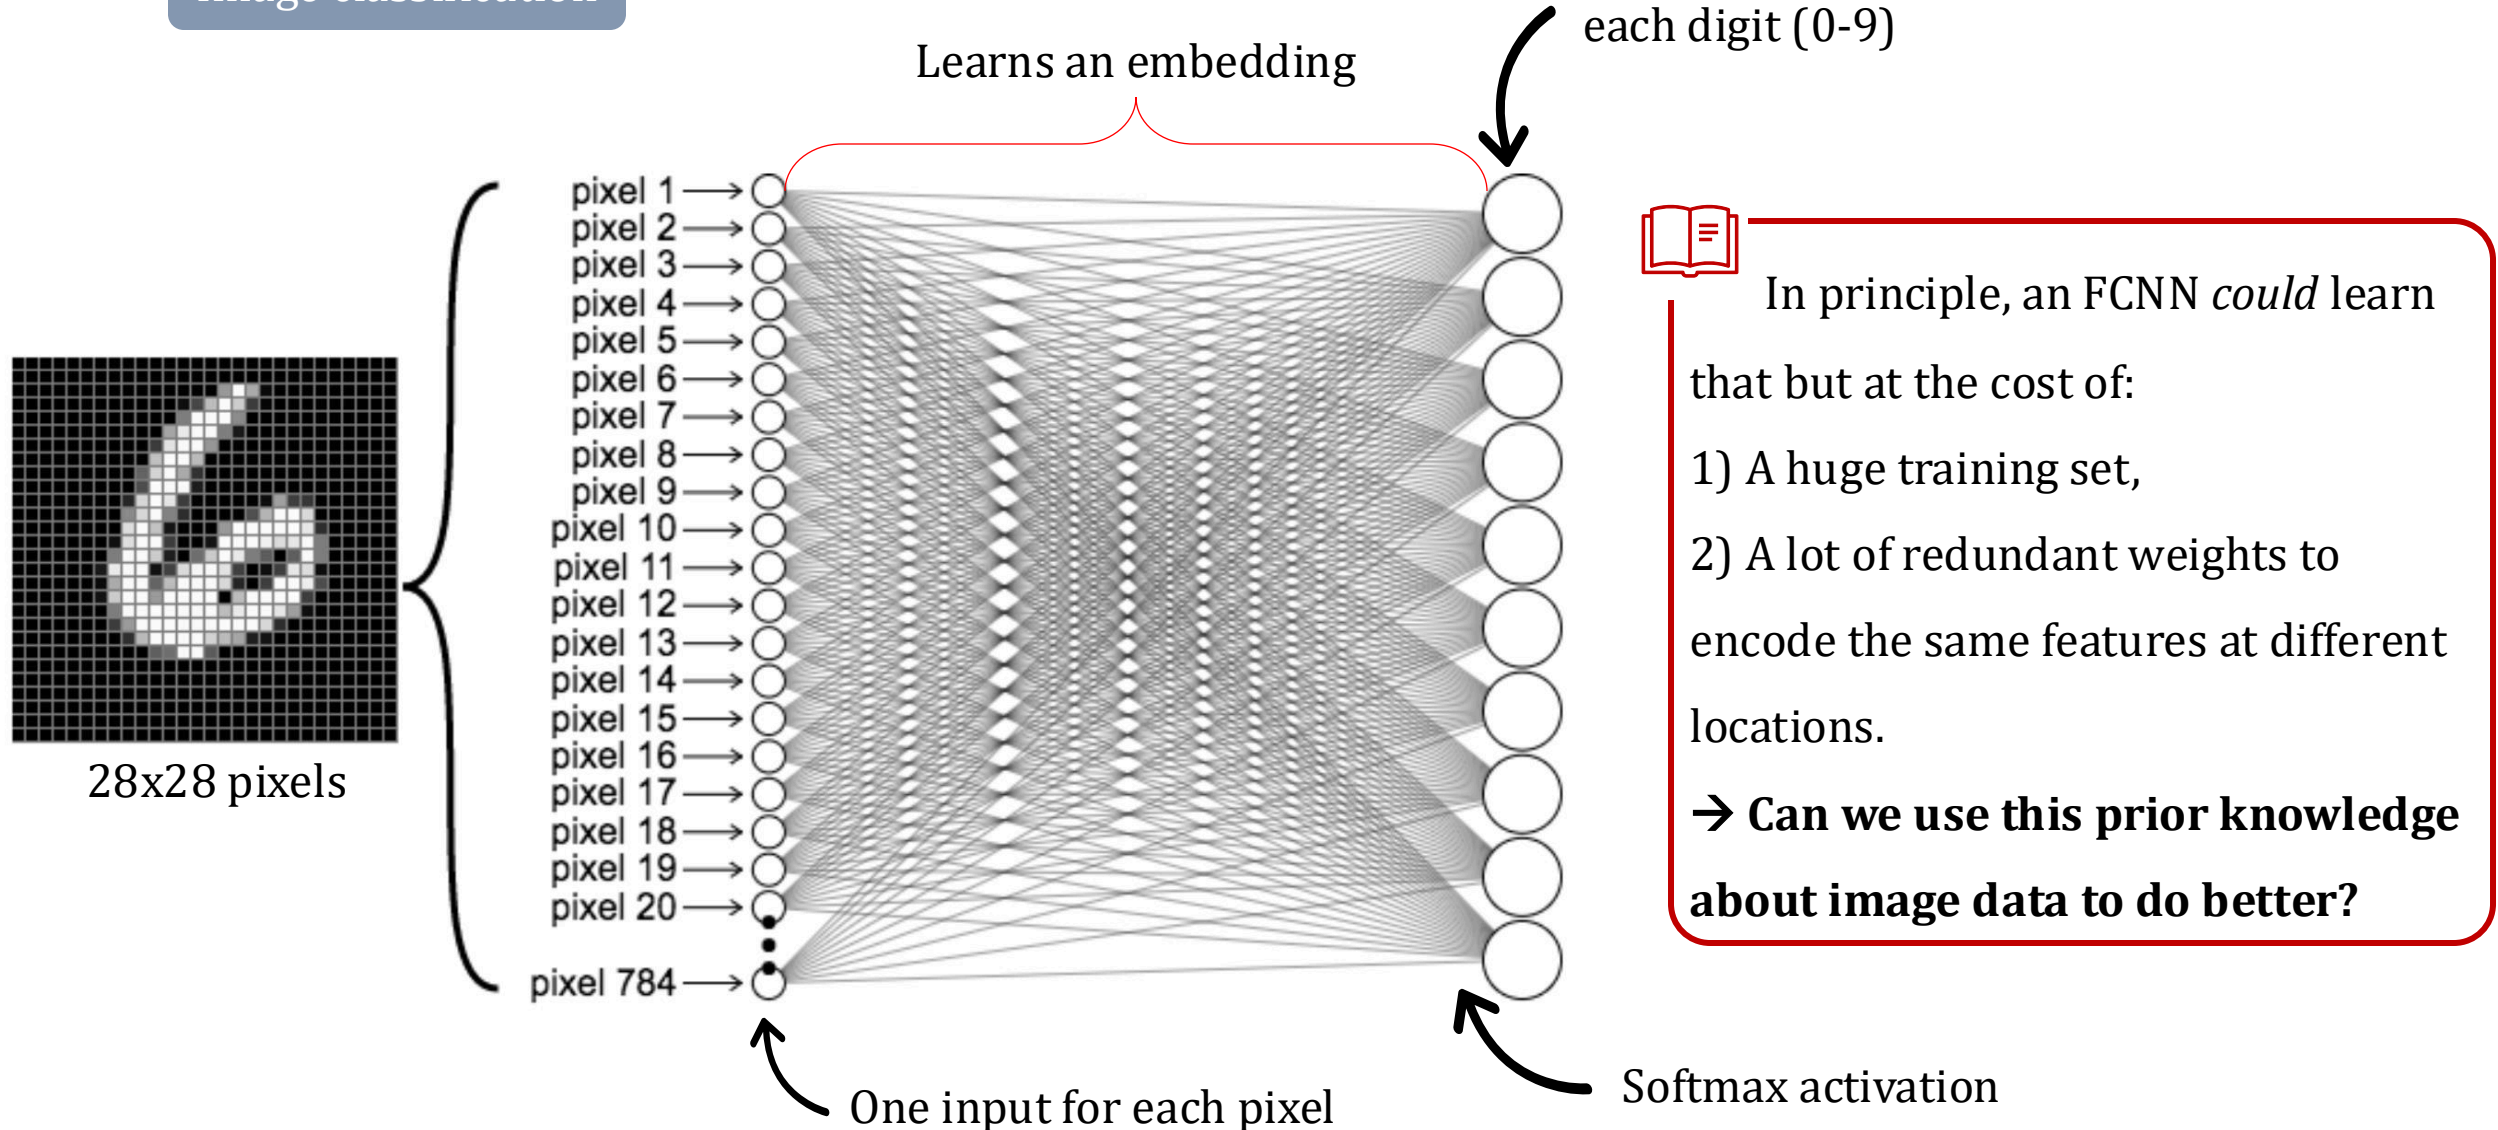
Why going beyond fully-connected neural network?

- Consider the simplest task with **classification**
- Standard images taken by a current smartphone are of size $4032 \times 3024 = 12\text{M}$ pixels
- A 2 layer-FCNN with just 100 (!) units has $> 1\text{B}$ parameters
- FCNN are **unstructured** with no invariance with respect to **translations** or **local distortions**



They are all zeros!

Image classification



Convolution

1D


$$(x * w)(t) = \sum_{k=-K}^K x(k)w(t - k)$$

2D

$$(\mathbf{X} * \mathbf{W})_{ij} = \sum_{l=1}^L \sum_{k=1}^K x_{lk} w_{i-l, j-k}$$

Convolution operation is a **linear operation**
equivariant to translation

$$\begin{aligned} \sum_{l=1}^L \sum_{k=1}^K x_{l+c, k+c} w_{i-l, j-k} &= \sum_{l=1}^L \sum_{k=1}^K x_{lk} w_{i-l+c, j-k+c} \\ &= (\mathbf{X} * \mathbf{W})_{i+c, j+c} \end{aligned}$$



Animation from [Wikipédia](#)

Convolution

Example of a kernel for vertical edge detection

 X

1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	0	0	0

 W


1	0	-1
1	0	-1
1	0	-1

*

=

 $X * W$

0	3	3	0
0	3	3	0
0	3	3	0
0	3	3	0



The border is enhanced!

Convolution

Example of a kernel for vertical edge detection

$$\begin{array}{c} \mathbf{X} \\ \begin{array}{|c|c|c|c|c|c|c|c|} \hline 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \hline 0 & 1 & 1 & 1 & 0 & 0 & 0 & 0 \\ \hline 0 & 1 & 1 & 1 & 0 & 0 & 0 & 0 \\ \hline 0 & 1 & 1 & 1 & 0 & 0 & 0 & 0 \\ \hline 0 & 1 & 1 & 1 & 0 & 0 & 0 & 0 \\ \hline 0 & 1 & 1 & 1 & 0 & 0 & 0 & 0 \\ \hline 0 & 1 & 1 & 1 & 0 & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \hline \end{array} \end{array} * \begin{array}{c} \mathbf{W} \\ \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline \end{array} \end{array} = \begin{array}{c} \mathbf{X * W} \\ \begin{array}{|c|c|c|c|c|c|} \hline -2 & 0 & 2 & 2 & 0 & 0 \\ \hline -2 & 0 & 3 & 3 & 0 & 0 \\ \hline -3 & 0 & 3 & 3 & 0 & 0 \\ \hline -3 & 0 & 3 & 3 & 0 & 0 \\ \hline -2 & 0 & 3 & 3 & 0 & 0 \\ \hline -2 & 0 & 2 & 2 & 0 & 0 \\ \hline \end{array} \end{array}$$

\mathbf{X} and $\mathbf{X * W}$ have the same size




In practice, we use **padding** to add zeros around the image \mathbf{X} so that the border of the image are seen as much as other pixels.

Convolutional layers

- A **convolutional layer** in a neural network is simply implementing the convolution operation with a filter (or kernel) **W shared across all the inputs**
- In 2 dimensions with a first layer of the same size as the input, a square weight matrix has $K \times K$ elements, while a fully connected layer has $d \times d$
- Typically, $K \ll d, K \sim O(1)$, usually 3 or 5.
- Units of a given layer ℓ only sees a subset of the activations of the previous layer $\ell - 1$: **local receptive field**
- **Deeper units** in the network **are influenced by more inputs**, hence learning higher-order features



The **convolutional layer** has three properties: **sparse interactions**, **parameters sharing** and **equivariant representation**.




Pooling layers

- To make the network **invariant to slight local translations**, we need an additional element
- Idea: local invariance to translations of features in a given window
- In 2D

$$z_{ij}^{[\ell+1]} = \max_{(r,s) \in K_1 \times K_2} z_{rs}^{[\ell]}$$

$$z_{ij}^{[\ell+1]} = \frac{1}{K_1 \times K_2} \sum_r^{K_1} \sum_s^{K_2} z_{rs}^{[\ell]}$$




A max non-overlapping pooling layer with window of size 2x2



The **pooling layer** grants the network **some local invariance to translation** and **reduces the image size**


CNNs

A convolutional neural network (CNN) is a cascade of **convolutional layers** and **pooling layers** to ensure **invariance to small (local) shifting, scaling and distortions** through local receptive field, shared weights and spatial sub-sampling.




Architecture of LeNet-5 used for digit recognition in [LeCun+98](#). It has 60,000 trainable parameters.

More specific, large-scale,
high-order, features



Layer 3



Layer 2

Local, first-order,
features (edges)



Layer 1

Depth

Image from [Albawi et al., 2017](#)